

Small-scale training loss curve (b=4, hidden=64, iters=200)

