

Quality vs inference time under compute budgets

